



Detecting Credit Card Fraud

Final Project - W207 - Section 11 - Fall 2022
Naikaj Pandya, Chase Madson, Eric Danforth

The Competition: *IEEE-CIS Fraud Detection*

Research Prediction Competition

IEEE-CIS Fraud Detection

Can you detect fraud from customer transactions?

\$20,000
Prize Money

IEEE Computational Intelligence Society • 6,351 teams • 3 years ago

[Overview](#) [Data](#) [Code](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#) [Submissions](#) [Late Submission](#) ...

Overview

Description

Evaluation

Prizes

Timeline

Imagine standing at the check-out counter at the grocery store with a long line behind you and the cashier not-so-quietly announces that your card has been declined. In this moment, you probably aren't thinking about the data science that determined your fate. Embarrassed, and certain you have the funds to cover everything needed for an epic nacho party for 50 of your closest friends, you try your card again. Same result. As you step aside and allow the cashier to tend to the next customer, you receive a text message from your bank. "Press 1 if you really tried to spend \$500 on cheddar cheese."

[Link to competition](#)

- Hosted by IEEE-CIS in 2019 to improve fraud prevention system
- ```
train.sample(5)
```

  
[10] ✓ 1.7s  
...

|        | TransactionID | isFraud | TransactionDT | TransactionAmt |
|--------|---------------|---------|---------------|----------------|
| 336447 | 3323447       | 0       | 8281233       | 77.95          |
| 443089 | 3430089       | 0       | 11273831      | 57.95          |
| 4917   | 2991917       | 0       | 170923        | 44.00          |
| 534610 | 3521610       | 0       | 14076224      | 226.00         |
| 9535   | 2996535       | 1       | 273968        | 100.00         |
- 590K rows × 430 features for training
  - One record per transaction
  - Binary target: **isFraud == 1**
- Our goal: Build a good NN classifier




# Available Features for Each Transaction

| Group of Features                            | Features                                                 | Example                                                                                       |
|----------------------------------------------|----------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| Basic Transaction Details                    | Timestamp, amount, payment card, email domains, addr1 &2 | Card level details (visa, MC)<br>Product CD (actual product/service)<br>Txn amount, date, etc |
| Card-Level Counts                            | C1-C14                                                   | number of addresses are associated with the payment card                                      |
| Card-Level Time Between (Unspecified) Events | D1-D15                                                   | days between prior transactions                                                               |
| Card-Level Matches                           | M1-M9                                                    | match, such as names on card and address, etc.                                                |
| Various Vesta-Enriched Features              | V1-V339                                                  | Vesta engineered rich features, including ranking, counting, and other entity relations.      |

# Dealing with Imbalanced Data

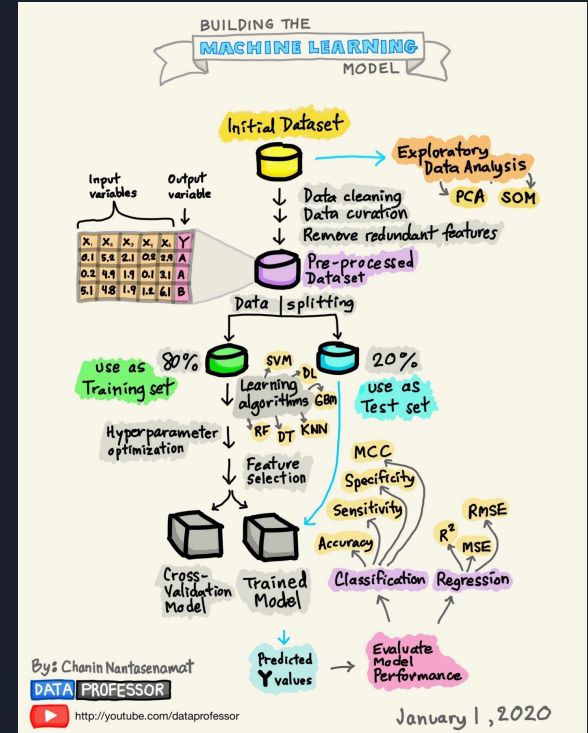
- Fraud detection training dataset is highly imbalanced.
- The training dataset consisted of 96.5% non-fraud transactions and 3.5% fraudulent transactions
- A simple-majority baseline would yield a 96.5% accurate model, but would not be too good at identifying fraud.
- Downsampling: Kept all frauds, and then down sampled to get the same number of random, non-fraud transactions



<https://datascience.aero/predicting-improbable-part-1-imbalanced-data-problem/>

# ML Model Design Process Overview

- The problem: to predicting whether a given transaction is fraudulent.
- EDA: Look at the feature space of more than 390 columns
- Baseline Model: KNN (5)
- Neural Network model:
  - Deeper EDA
  - Data processing, cleanup
  - Feature engineering (iterative)
  - Build & fit model




# Baseline Model

- Simple majority baseline model would not work due to heavy data imbalance.
- We decided on using KNN model, since we theorized that fraudulent transactions could have patterns and clustering in certain features.
- 18 features to develop KNN model: TransactionDT, TransactionAmt, ProductCD, card1, and C1-C14
- Baseline model has validation AUC of 86.9% and a Kaggle score of 82.5%

## KNN Baseline Score

Public: 86.9%

Private: 82.5%



<https://www.datacamp.com/tutorial/k-nearest-neighbor-classification-scikit-learn>

# Data Pre-Processing & Feature Selection

De-anonymization

Dimensionality reduction

- Limited PCA given time available and number of anonymized features

- Dropped identity columns in the transaction table that weren't providing much information

- Identity table fields such as 'browser' and 'OS' mapped down to family groups

Min-max normalization


Missing values replaced for numeric values

Categorical features one-hot encoded

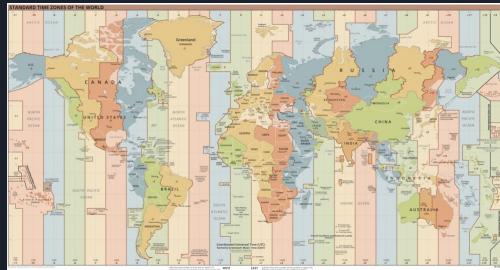


# Feature Engineering

- Time & location features - training data range
- Top level domains for recipient and purchaser emails
- Transaction in foreign currency - 3rd decimal place
- Create heuristic for User ID
- Filled identity table & De-anonymized features
  - Country - Most fraud originating from Russia and China
  - Timezone
- Foreign currency conversion



<https://www.xe.com/currencycharts/?from=USD&to=EUR&view=5Y>



[https://en.wikipedia.org/wiki/File:World\\_Time\\_Zones\\_Map.png](https://en.wikipedia.org/wiki/File:World_Time_Zones_Map.png)



# Keras Model


- Final model used Keras Sequential API
- 4 Dense layers, each with Batch Normalization and Dropout applied
- Total parameters: 730,881
- We had a working functional model as well
- This gave us flexibility to conduct data preprocessing within the model
- We found that some Tensorflow features are hardware specific and not easily portable to accelerators



| Layer (type)                       | Output Shape | Param # |
|------------------------------------|--------------|---------|
| dense_0 (Dense)                    | (None, 512)  | 496128  |
| batch_norm_0 (Batch Normalization) | (None, 512)  | 2048    |
| dropout_0 (Dropout)                | (None, 512)  | 0       |
| dense_1 (Dense)                    | (None, 256)  | 131328  |
| batch_norm_1 (Batch Normalization) | (None, 256)  | 1024    |
| dropout_1 (Dropout)                | (None, 256)  | 0       |
| dense_2 (Dense)                    | (None, 256)  | 65792   |
| batch_norm_2 (Batch Normalization) | (None, 256)  | 1024    |
| dropout_2 (Dropout)                | (None, 256)  | 0       |
| dense_3 (Dense)                    | (None, 128)  | 32896   |
| batch_norm_3 (Batch Normalization) | (None, 128)  | 512     |
| dropout_3 (Dropout)                | (None, 128)  | 0       |
| dense (Dense)                      | (None, 1)    | 129     |
| Total params: 730,881              |              |         |
| Trainable params: 728,577          |              |         |
| Non-trainable params: 2,304        |              |         |

# Accelerators - TPUs

- We experimented with accelerators to speed up training time to allow more experimentation and larger models
- Training epochs were lightning fast even with thousands of features; however, model serialization proved to be a memory intensive operation
- Overall training time for a model of this size was not improved by using TPUs
- Performance metrics were degraded compared to the same model architecture on a CPU, possibly due to a reduction in decimal precision


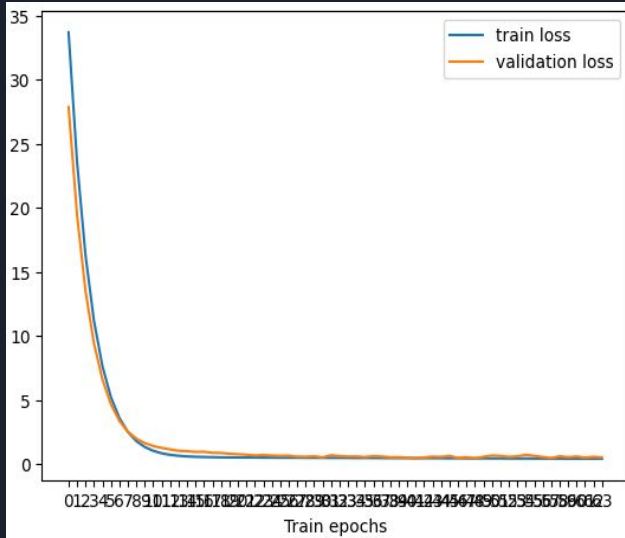


<https://blog.google/technology/developers/io21-helpful-google/>

# Hyperparameter Tuning

| Hyperparameter Type          | Hyperparameter                     | Selection                                        |
|------------------------------|------------------------------------|--------------------------------------------------|
| Training and backpropagation | Loss Function and Validation Split | Binary cross-entropy w/ 10% split                |
|                              | Optimizer and Learning Rate        | Adam w/ L.R. = 0.0005                            |
|                              | Epochs and Batch Size              | 64 epochs in batches of 1024                     |
| Hidden Layers                | Hidden Layers and Nodes            | [512, 256, 256, 128]                             |
|                              | Activation Function                | ReLU                                             |
| Regularization               | Model Parameters                   | L2 regularization w/ lambda = 0.05               |
|                              | Hidden Layers                      | Batch normalization and dropout at a rate of 50% |

# Model Performance and Results



Best Kaggle Score

Public: 87.3%

Private: 84.9%

KNN Baseline Score

Public: 86.9%

Private: 82.5%

# Data Generating Process for our Labels

What we expected




<https://stressexterminators.wordpress.com/2015/02/10/stress-hacking-chris-tmas-lights-organization/>

What we got



# Train New Model on Latest Transaction by Card

| card1 | addr1 | date_account_opened | P_emaildomain | isFraud | txn_timestamp             |
|-------|-------|---------------------|---------------|---------|---------------------------|
| 7505  | 433.0 | 2018-04-14          | yahoo.com     | 1       | 2018-08-13 15:22:46-07:00 |
| 7505  | 433.0 | 2018-04-14          | yahoo.com     | 1       | 2018-08-20 02:12:01-07:00 |
| 7505  | 433.0 | 2018-04-14          | yahoo.com     | 1       | 2018-09-02 16:06:44-07:00 |
| 7585  | 433.0 | 2018-04-14          | yahoo.com     | 1       | 2018-09-06 18:48:05-07:00 |
| 7505  | 433.0 | 2018-04-14          | yahoo.com     | 1       | 2018-10-21 14:54:18-07:00 |
| 7585  | 433.0 | 2018-04-14          | yahoo.com     | 1       | 2018-11-25 01:00:25-08:00 |



Unshrink

Kaggle Score  
Public: 85.9%  
Private: 84.3%  
Val. AUC: 85.6%

# A Very Simple Ensembling Approach

Best Model

Kaggle Score  
Public: 87.3%  
Private: 84.9%

Using Latest  
Transaction

Kaggle Score  
Public: 85.9%  
Private: 84.3%

Simple Average  
of the Two

Kaggle Score  
Public: 88.4%  
Private: 86.6%





# Things we Tried, or Would Try in Future

## Trial and Error


- ❖ PCA dim. reduction of the 339 Vesta Features
- ❖ Impute median on missing values
- ❖ Log transform heavy-skewed features

## Looking ahead

- ❑ More feature engineering
- ❑ Use just the “top 20” Vesta Features
- ❑ Use more data from the identity table provided
- ❑ Better approximation to identify genuine fraud and/or an ID for card
- ❑ Smarter missing value imputation
- ❑ Categorical embeddings - cat2vec



# Questions?



<https://developer.confluent.io/tutorials/credit-card-activity/confluent.html>



Thank You!